

ORACLE®

Oracle数据库性能调优

金丹

资深技术顾问



最重要的规则

• 最少的参数,最少的性能影响

• 除非你很清楚你在干什么,否则不要设置_****

· 除非你很清楚你在干什么,否则不要设置 event

定期监控

- DB statspack,AWR,ADDM
- OS ,vmstat,sar....
- Network
- Grid Control
- RUEI

Oracle Tuning Methods

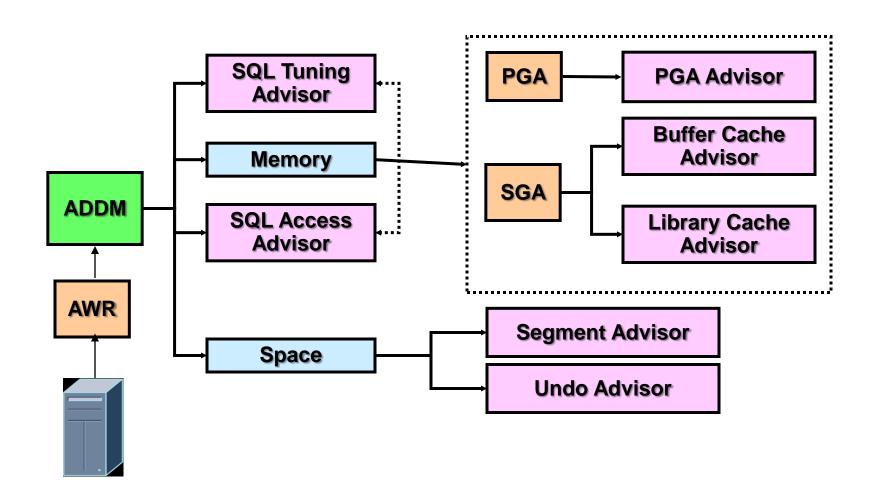
- Prehistoric (v5)
 - Debug code
- Dark Ages (v6)
 - Counters/Ratios
 - BSTAT/ESTAT
 - SQL*Trace

- Renaissance (v7)
 - Introduction of Wait Event instrumentation
 - Move from counters to timers
 - STATSPACK

The Modern Age of Time-based Tuning

- YAPP (8i) Instance tuning using instance statistics
 - Non intrusive, always on
 - Broadly scoped
- Method R (9i) Session tuning using 10046 SQL traces
 - Tightly scoped
 - Must be highly selective
- DB Time Tuning (10g) Comprehensive tuning using fundamental notion of time in database
 - Multiple scoping levels
 - Always on, non-intrusive
 - Built into infrastructure: instrumentation, ASH, AWR, ADDM, EM

Oracle 10g/11g的智能化工具

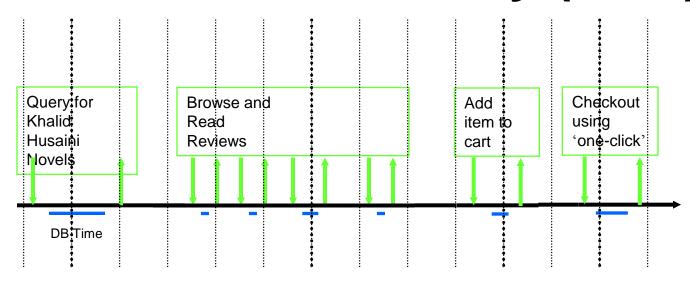


Active Session History (ASH)

- Samples active sessions every one second into memory (v\$active_session_history)
- Direct access to kernel structures
- One of ten samples flushed to AWR at every snapshot
- Data captured includes:
 - SID
 - SQL ID
 - Program, Module, Action
 - Wait event#
 - Object, File, Block
 - actual wait time (if captured while waiting)

Enables targeted performance analysis of transient problems

Active Session History (ASH)



Time	SID	Module	SQL ID	State	Event
7:38:26	213	Book by author	qa324jffritcf	WAITING	db file sequential read
7:42:35	213	Get review id	aferv5desfzs5	CPU	
7:50:59	213	Add to cart	hk32pekfcbdfr	WAITING	buffer busy wait
7:52:33	213	One click	abngldf95f4de	WAITING	log file sync

ASH Report: Main Sections

ASH Report For BUGAP/bug1ap

(1 Report Target Specified)

DB Name	DB ld	Instance	Inst num	Release	RAC	Host
BUGAP	1679034986	bug1ap	1	10.2.0.1.0	YES	dbs232

CPUs	SGA Size	Buffer Cache	Shared Pool	ASH Buffer Size
4	2,576M (100%)	1,200M (46.6%)	1,109M (43.0%)	8.0M (0.3%)

	Sample Time	Data Source		
Analysis Begin Time:	21-Sep-06 13:13:20	V\$ACTIVE_SESSION_HISTORY		
Analysis End Time:	21-Sep-06 13:18:20	V\$ACTIVE_SESSION_HISTORY		
Elapsed Time:	5.0 (mins)	Missing 1.0 mins (20%) of activity		
Sample Count:	1,330			
Average Active Sessions:	4.43			
Avg. Active Session per CPU:	1.11			
Report Target:	SQL_ID like 'cyaj7dkrbqs95'	4% of total database activity		



- Top Events
- Load Profile
- Top SQL
- Top Sessions
- Top Objects/Files/Latches
- Activity Over Time





ASH Report: Top Events

Top User Events

Event	Event Class	% Activity	Avg Active Sessions
db file sequential read	User I/O	68.80	3.05
gc buffer busy	Cluster	12.33	0.55
buffer busy waits	Concurrency	9.25	0.41
read by other session	User I/O	5.64	0.25
gc cr disk read	Cluster	1.28	0.06

ASH Report: Top SQL

Top SQL Statements										
xt	SQL Text	% Event	Event	% Activity	Planhash	SQL ID				
ct_id, a.t	select a.product_id,	67.67	db file sequential read	97.74	2770286798	cyaj7dkrbqs95				
		12.18	gc buffer busy							
		9.02	buffer busy waits							

Complete List of SQL Text SQL Text SQL Id cyaj7dkrbqs95 | select a.product_id, a.test_name, a.rptnol, a.status, to_char(a.rptdate, 'mm/dd/yyyy'), a.base_rptno, b.product_id, a.do_by_release from rpthead a, rpthead b where a.product_id in (501) and a.test_name is not NULL and (a.subject like '%R12.FIN.A.XP%' or a subject like '%R12.FIN.A.QA%' or a subject like '%R12.FIN.A.UI%') and a base_rptno = b.rptno(+) order by a.test_name, a.rptno

ASH Report: Activity Over Time

Activity Over Time

- · Analysis period is divided into smaller time slots
- . Top 3 events are reported in each of those slots
- · 'Slot Count' shows the number of ASH samples in that slot
- · 'Event Count' shows the number of ASH samples waiting for that event in that slot
- . '% Event' is 'Event Count' over all ASH samples in the analysis period

Slot Time (Duration)	Slot Count	Event	Event Count	% Event
13:14:00 (1.0 min)	220	db file sequential read	163	12.26
		gc buffer busy	27	2.03
		buffer busy waits	12	0.90
13:15:00 (1.0 min)	295	db file sequential read	222	16.69
		gc buffer busy	32	2.41
		buffer busy waits	22	1.65
13:16:00 (1.0 min)	305	db file sequential read	211	15.86
		gc buffer busy	43	3.23
		read by other session	23	1.73
13:17:00 (1.0 min)	295	db file sequential read	199	14.96
		gc buffer busy	35	2.63
		buffer busy waits	28	2.11
13:18:00 (20 secs)	100	db file sequential read	46	3.46
		buffer busy waits	27	2.03
		gc buffer busy	14	1.05

Automatic Workload Repository (AWR)

- Built-in workload and performance statistics repository in the database
- Automatically Captures Workload Data
 - Every 60 minutes, or manually, saves data for 8 days by default
- Resides in SYSAUX tablespace
- Space requirements automatically managed
 - Old data is automatically purged nightly based on retention interval
- Stores different classes of data:
 - BASE STATISTICS e.g., physical reads
 - SQL STATISTICS e.g., disk reads (per sql stmt)
 - METRICS e.g., physical reads / sec
 - ACTIVE SESSION HISTORY (ASH)



AWR Compare Period Report

WORKLOAD REPOSITORY COMPARE PERIOD REPORT

Snapshot Set	DB Name	DB Id	Instance	Inst num	Release	Cluster	Host	Std Block Size
First (1st)	PROD	1545480911	prod	1	11.1.0.6.0	NO	tdsrat01 - d1 .oracleads.com	8192
Second (2nd)	DEV	3707267492	dev	1	11.1.0.6.0	NO	tdsrat01- d2.oracleads.com	8192

Snapshot Set	Begin Snap Id	Begin Snap Time	End Snap Id		Avg Active Users		DB time (min)
1st	339	27-Mar-08 20:17:34 (Thu)	340	27-Mar-08 20:22:07 (Thu)	0.68	4.54	3.09
2nd	396	27-Mar-08 20:52:52 (Thu)	397	27-Mar-08 20:56:18 (Thu)	0.30	3.44	1.04
%Diff					-55.88	-24.23	-66.40

AWR Compare Period Report: Load Profile

Load Profile						
	1st per sec	2nd per sec	%Diff	1st per txn	2nd per txn	%Diff
DB time:	0.68	0.30	-55.88	0.09	0.03	-66.67
CPU time:	0.36	0.20	-44.44	0.05	0.02	-60.00
Redo size:	141,784.30	186,369.33	31.45	18,478.25	18,542.89	0.35
Logical reads:	30,539.38	1,289.19	-95.78	3,980.09	128.27	-96.78
Block changes:	726.20	949.25	30.71	94.64	94.45	-0.20
Physical reads:	6,790.88	0.61	-99.99	885.03	0.06	-99.99
Physical writes:	2.88	1.68	-41.67	0.38	0.17	-55.26
User calls:	338.11	447.90	32.47	44.06	44.56	1.13
Parses:	15.58	17.39	11.62	2.03	1.73	-14.78
Hard parses:	0.83	0.24	-71.08	0.11	0.02	-81.82
Sorts:	4.57	9.52	108.32	0.60	0.95	58.33
Logons:	0.09	0.11	22.22	0.01	0.01	0.00
Executes:	344.89	449.90	30.45	44.95	44.76	-0.42
Transactions:	7.67	10.05	31.03			
				1st	2nd D	iff
% Blocks changed	per Read:			2.38	73.63	71.25
Recursive Call %:				28.01	21.66	-6.35
Rollback per transa	action %:			0.96	0.68	-0.28
Rows per Sort:				51.54	11.15	-40.39
Avg DB time per Ca	all (sec):	0.00	0.00	-0.00		

AWR Compare Period Report: Top Timed Events

Top Timed Events

 Events with a "-" did not make the Top list in this set of snapshots, but are displayed for comparison purposes

1st						2nd					
Event	Wait Class	Waits	Time (s)	Avg Time (ms)	%DB time	Event	Wait Class	Waits	Time (s)	Avg Time (ms)	%DB time
CPU time			97.03		52.28	CPU time			40.26		64.57
db file sequential read	User I/O	1,218	11.76	9.66	6.34	log file parallel write	System I/O	1,854	5.83	3.14	9.35
log file parallel write	System I/O	2,382	7.43	3.12	4.00	SQL*Net message to client	Network	91,222	1.50	0.02	2.40
log file sync	Commit	2,018	7.21	3.57	3.88	latch: shared pool	Concurrency	3	0.62	206.52	0.99
SQL*Net message to client	Network	91,076	1.56	0.02	0.84	os thread startup	Concurrency	3	0.43	142.49	0.69
-os thread startup	Concurrency	6	1.18	197.29	0.64	-log file sync	Commit	35	0.27	7.66	0.43
-						-db file sequential read	User I/O	118	0.11	0.96	0.18

数据库性能基础

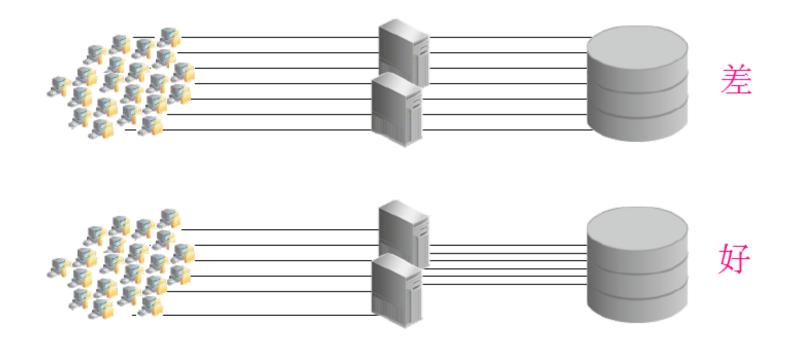
- 好的Schema和数据模型设计是好的数据库性能的基础。
- 好的Schema和数据模型设计应该使应用操作更加简单 并减少错误发生
 - 选择正确的数据类型(e.g. Char vs VARCAHR2)
 - 使用不同的索引设计高效率的访问路径
 - 选择正确的存储选项
 - 使用分区
 - 收集并维护准确的统计信息
 - 编写简单高效的SQL
 - 尽量避免序列化瓶颈

数据库性能基础

- 应用服务器/中间件管理数据库会话连接,保证数据库的性能和稳定
 - 在一台机器上超过5000进程会使机器变的不稳定
 - 在实验室很容易模拟大量的数据库连接,但是在生产系统上要使用则非常困难
 - logon/logoff对性能的影响进程的建立和撤销开销非常大
 - 连接和去连接到共享内存的开销也非常大
- 共享SQL, 在OLTP应该中使用绑定变量
- 应用应该尽量避免对数据库对象进行creation、drop 、truncate等操作



会话管理



- 不要连接所有浏览器用户
- 不要连接、执行操作、断开连接
- 使用连接池—中间层的最大优点

性能与数据安全的一些取舍

- 影响性能的一些数据安全设置:
 - 多镜像控制文件
 - 重做日志组中多镜像日志组成员
 - 频繁Check Point
 - 备份数据文件
 - 归档状态
 - 块校验
 - 并发用户访问和事务

数据库性能基础

CPU and Memory

- For OLTP workloads the CPU utilization should not exceed 65-70%
- Scale Nodes/CPUs/Cores as a power of 2.
- Start with 4Gig Memory per Core and expand according to workload type

1/0

- Target 5-10 millisec response time for disks performing response time critical I/O.
- Start by assuming 30 IOPS per disk for OLTP and 20MBytes per Disk in DSS. This is way below theoretical values but allows for media repair etc.
- You can run HBAs and SAN switches at about 70-90% of theoretical values.
- Allocate and Administer Disk by bandwidth in addition to storage requirements.
- 200MB/s per CPU core

Interconnect

- For Data passing intensive DSS Clusters Balance equalize Interconnect bandwidth with I/O bandwidth. Complex queries running across a cluster will perform sorts and joins across the interconnect.
- In the future we may require higher interconnect bandwidth for DW/BI workloads because the data on disk is compressed
- Do not confuse bits for Bytes when capacity planning Interconnect or SAN.

Oracle的性能解决方案

- 绑定变量
- 并行查询
- 索引
- 物化视图
- 分区
- Real Application Cluster
- 压缩
- 结果集缓存
- Active Data Guard
- TimesTen

并行查询

- SELECT /*+ PARALLEL (SALES, 9) */ * FROM SALES;
- SELECT /*+ PARALLEL INDEX(c,ic,3)*/ * FROM customers c WHERE cust city = 'MARSEILLE';

并行DML

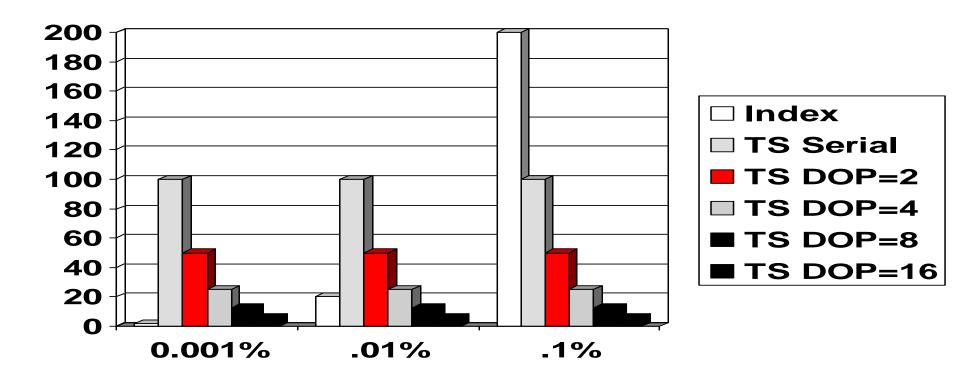
- ALTER SESSION ENABLE/DISABLE PARALLEL DML
- MERGE /*+ PARALLEL(c,3) PARALLEL(d,3) */ INTO customers c USING diff customers d ON (d.cust id = c.cust id) WHEN MATCHED THEN UPDATE SET c.cust last name = d.cust last name, c.cust city = d.cust city WHEN NOT MATCHED THEN INSERT (c.cust id, c.cust last name) VALUES (d.cust id, d.cust last name);

并行DDL

- CREATE INDEX
- CREATE TABLE ... AS SELECT
- ALTER INDEX ... REBUILD

并行查询性能分析

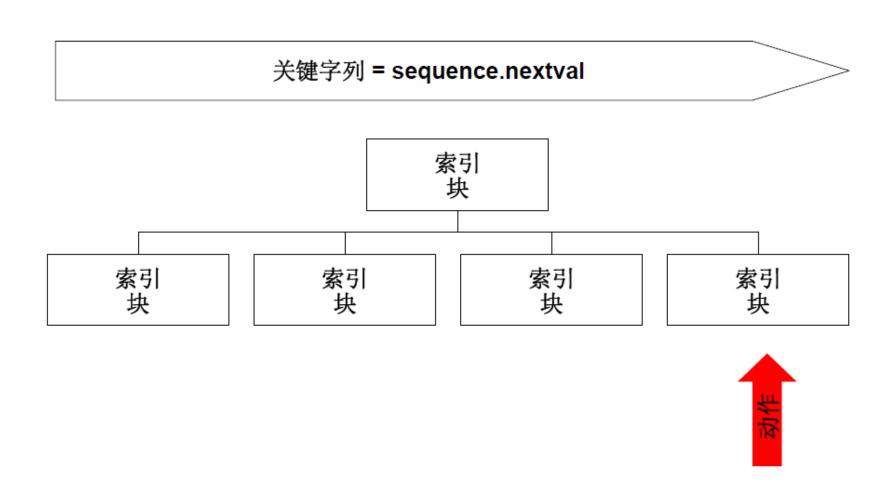
• 假设某张表拥有100,000,000 行记录数,每行记录100 Byte 大小,表扫描I/O速率是每秒1M/10ms,索引I/O是8k/5ms并且假设缓存命中率为80%



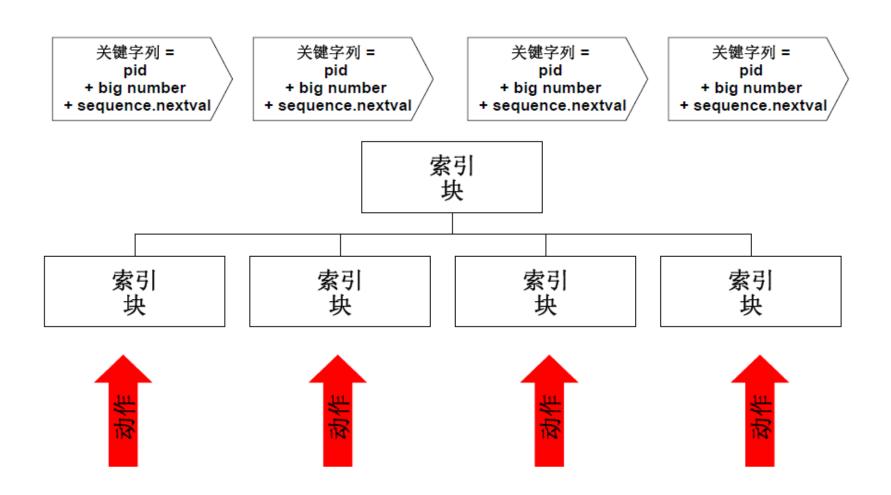
索引

- 索引的影响
 - 加快能适当使用索引的SELECT, UPDATE
 - 减慢INSERT速度
 - 需要额外空间
- 索引使用原则
 - 只有需要的时候才创建相应的索引
 - 为解决某个问题创建一个索引可能会影响其他的SQL执行
 - 删除没用的索引

使用索引:集中IO



使用索引:分散IO



索引的类型

- 唯一及非唯一索引
- 复合索引
- 按索引存储方法
 - B*-tree
 - 普通
 - 反键索引
 - 函数索引
 - 位图索引

• 索引组织表

物化视图

- 物化视图:
 - 是一个预计算的结果集
 - 拥有自己的数据存储:
 - 拥有自己的索引
 - 适用场合:
 - 复杂的join
 - 概要和聚集数据集

分区的作用

- 分区对性能影响
 - 分区忽略
 - 并行处理

数据分区技术大大提高访问速度 — "有的放矢、并 行处理"

select sum(amount)

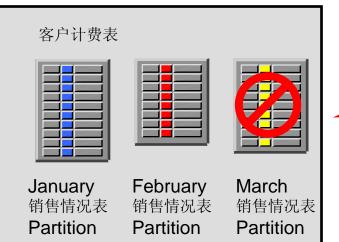
销售情况表 From Where 日期

and

between 20-JAN-2004 5-FEB-2004



- 分区可以显著提高访问大表 时的性能
- 分区的存在对应用系统是透 明的

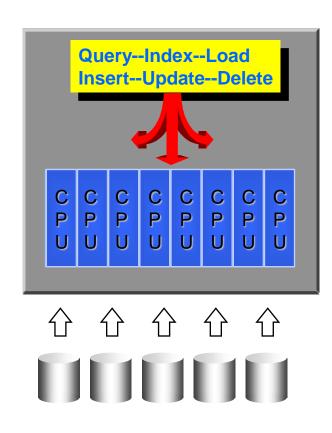


分区忽略技术

系统核心自动根据分区情况优 化数据访问, 忽略无关的数据 分区



强大的并行处理能力



- 随着数据库的增长,必须采用 并行处理方式以保证响应时间
- 在数据库中进行并行处理可以 显著提高批量操作的性能
 - 批量更新
 - 批量删除
 - 批量插入
- 并行执行 SQL语句
- 数据动态分片
- RAC的多节点并行更新



EXPLAIN PLAN命令: 示例1

```
EXPLAIN PLAN

SET STATEMENT_ID = 'demo01' FOR

SELECT e.last_name, d.department_name

FROM hr.employees e, hr.departments d

WHERE e.department_id = d.department_id;

Explained.
```

EXPLAIN PLAN 命令并不实际执行该SQL。



EXPLAIN PLAN命令: 示例2

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY());
```

```
Plan hash value: 2933537672
                        | Name | Rows | Bytes | Cost (%CPU|
| Id | Operation
                                        | 106 | 2862 | 6 (17|
   0 | SELECT STATEMENT
                    | | 106 | 2862 | 6 (17)
 1 | MERGE JOIN
 2 | TABLE ACCESS BY INDEX ROWID| DEPARTMENTS | 27 | 432 | 2 (0|
 3 | INDEX FULL SCAN | DEPT_ID_PK | 27 | | 1 (0|

4 | SORT JOIN | 107 | 1177 | 4 (25|

5 | TABLE ACCESS FULL | EMPLOYEES | 107 | 1177 | 3 (0|
|* 4 | SORT JOIN
Predicate Information (identified by operation id):
  4 - access("E"."DEPARTMENT ID"="D"."DEPARTMENT ID")
      filter("E"."DEPARTMENT ID"="D"."DEPARTMENT ID")
18 rows selected.
```

SQL*Plus中设置AUTOTRACE: 示例1

• 启动 AUTOTRACE

set autotrace on

• 隐藏输出

set autotrace traceonly

- 只显示执行计划
 - AUTOTRACE需要实际执行SQL语句

set autotrace traceonly explain

SQL*Plus中设置AUTOTRACE: 示例2

```
set autotrace traceonly statistics
SELECT *
FROM products;
```

```
1 recursive calls
0 db block gets
9 consistent gets
3 physical reads
0 redo size
15028 bytes sent via SQL*Net to client
556 bytes received via SQL*Net from client
6 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
72 rows processed
```

哪些因素影响执行计划

- 统计信息
 - 对象统计信息
 - 表
 - 列
 - 索引
 - 系统统计信息
 - I/O 性能
 - CPU 性能
- 优化器、Hints
- 初始化参数
- 索引
- 分区
- 物化视图
- 并行
-

统计信息如何收集

- 自动收集统计信息
 - GATHER STATS JOB
- 手动收集统计信息
 - 使用DBMS STATS包
- 动态采样
 - optimizer_dynamic_sampling

优化器

- RBO
- CBO
 - ALL_ROWS
 - FIRST_ROWS
 - FIRST_ROWS_n

Hints使用举例

```
UPDATE /*+ INDEX(p PRODUCTS PROD CAT IX)*/
products p
SET p.prod min price =
        (SELECT
         (pr.prod list price*.95)
          FROM products pr
          WHERE p.prod id = pr.prod id)
WHERE p.prod category = 'Men'
AND p.prod status = 'available, on stock'
```

Hints类型

- 优化器模式
 - ALL_ROWS
- 访问路径
 - FULL
- 查询转换
 - STAR TRANSFORMATION
- Join顺序
 - ORDERED
- Join操作
 - USE_NL, USE_HASH
- 其他
 - APPEND

SQL Trace

• 针对当前会话:

```
SQL> ALTER SESSION SET sql_trace = true;
```

• 针对任意会话:

```
SQL> EXECUTE dbms_session.set_sql_trace(true);
```

```
SQL> EXECUTE dbms_system.set_sql_trace_in_session
2 (session_id, serial_id, true);
```

• 实例级,可以直接设置初始化参数:

```
SQL_TRACE = TRUE
```



TKPROF格式化trace文件

• TKPROF 命令示例:

```
OS> tkprof tracefile outputfile [options]
```

```
OS> tkprof
OS> tkprof ora_902.trc run1.txt
OS> tkprof ora_902.trc run2.txt sys=no sort=execpu print=3
```

TKPROF 格式化后输出:示例

```
select max(cust credit limit)
from customers
where cust city = 'Paris'
call
                         cpu elapsed disk
           count
                                                                query
                                                                           current
                                                                                          rows

      Parse
      1
      0.02
      0.02
      0
      0

      Execute
      1
      0.00
      0.00
      0
      0

      Fetch
      2
      0.10
      0.09
      1408
      1459

total 4 0.12 0.11 1408 1459
                                                                                  0
Misses in library cache during parse: 1
Optimizer mode: ALL ROWS
Parsing user id: 61
           Row Source Operation
Rows
       1 SORT AGGREGATE (cr=1459 pr=1408 pw=0 time=93463 us)
      77 TABLE ACCESS FULL CUSTOMERS (cr=1459 pr=1408 pw=0 time=31483 us)
```

10046事件分析

SQL> alter session set events '10046 trace name context forever,

level 8';

Session altered.

SQL> alter session set events '10046 trace name context off';

Session altered.

执行跟踪:

SQL> exec dbms_system.set_ev(5,323,10046,8, \scott');

PL/SQL procedure successfully completed.

结束跟踪:

SQL> exec dbms_system.set_ev(5,323,10046,0, \scott');

PL/SQL procedure successfully completed.

10053事件分析

SQL> alter session set events '10053 trace name context forever,

level 1';

Session altered.

SQL> alter session set events '10053 trace name context off';

Session altered.

执行跟踪:

SQL> exec dbms_system.set_ev(5,323,10053,1, \scott');

PL/SQL procedure successfully completed.

结束跟踪:

SQL> exec dbms_system.set_ev(5,323,10053,0, \scott');

PL/SQL procedure successfully completed.



ORACLE® 母母子